

Chapter 03: Filtering

Probabilistic Artificial Intelligence

Notes by Kumar Anurag

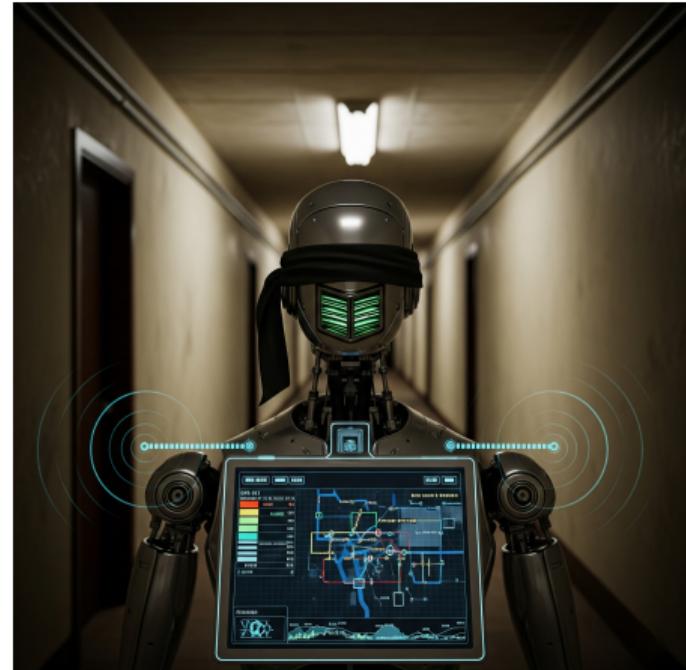
The Story of Blindfolded Robot

Imagine:

- ▶ A robot is navigating down a hallway, but it's **blindfolded**.
- ▶ It can't see where it is — it only has a noisy distance sensor and a map of the hallway.
- ▶ At each step, it must decide:
 - Where it *thinks* it is,
 - How to move next,
 - And how to adjust based on new noisy sensor readings.

The robot's challenge is exactly what filtering solves:

"Estimate where you are over time — with incomplete and uncertain information."



What is Filtering? The Core Problem

Filtering is about estimating the **hidden state** of a system over time, using a stream of noisy observations.

Why it's harder than regular regression

- ▶ The unknown (state) keeps changing.
- ▶ The observations are noisy and only partially informative.
- ▶ We must update our estimate *as new data arrives*.

We need:

- ▶ A **dynamics model**: How the state evolves over time.
- ▶ An **observation model**: How observations relate to the state.

Applications:

- ▶ Self-driving cars, speech recognition, finance, robotics, weather prediction...

State Space Models (SSMs)

Let's formalize the blindfolded robot's world.

- ▶ **Hidden States:** $X_t \in \mathbb{R}^d$
 - The robot's true position, speed, or internal state at time t .
 - Not directly visible — we want to estimate it!
- ▶ **Observations:** $Y_t \in \mathbb{R}^m$
 - Sensor readings at time t .
 - Noisy, partial glimpses of the hidden state.
- ▶ **Goal:** Given observations up to time t , i.e. $Y_{1:t}$,
 - Estimate the current state X_t .
 - Sometimes even forecast future states X_{t+1}, X_{t+2}, \dots

Key Idea

SSMs capture how the world evolves (*state transition*) and how we perceive it (*observation model*).

Bayesian Filtering: The Big Picture

Our Goal: Maintain a belief about the hidden state X_t as we receive observations $Y_{1:t}$.

Recursive Bayesian Estimation

At every time step, we perform two key operations:

1. Prediction (Time Update):

$$p(X_t | Y_{1:t-1}) = \int p(X_t | X_{t-1}) p(X_{t-1} | Y_{1:t-1}) dX_{t-1}$$

2. Update (Measurement Correction):

$$p(X_t | Y_{1:t}) \propto p(Y_t | X_t) p(X_t | Y_{1:t-1})$$

Repeat: These steps are performed in a loop as new data arrives.

Figure 3.1: How Bayesian Filtering Works

Explanation:

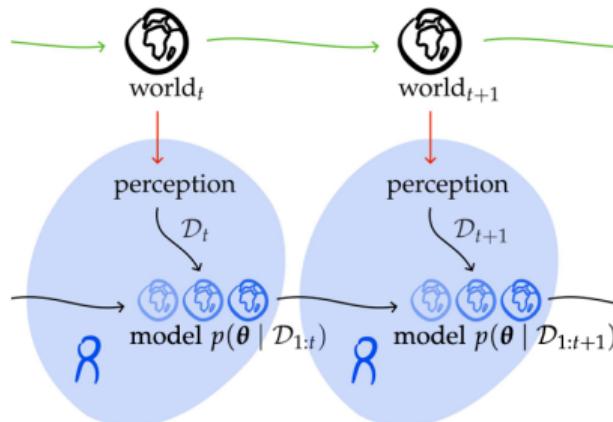
- The **true world state** evolves over time: $\text{world}_t \rightarrow \text{world}_{t+1}$.
- The agent perceives each state through noisy data: D_t, D_{t+1} .
- Based on its observations $D_{1:t}$, it maintains a belief (posterior) over the hidden variables:

$$p(\theta | D_{1:t})$$

- At each time step, the agent:
 1. **Updates** its belief using new data.
 2. **Predicts** the next state of the world.

Core Idea

Filtering is the cycle of *observe* → *update belief* → *predict* → *repeat*.



Introducing the Kalman Filter

Kalman Filter: A powerful, tractable solution to the filtering problem — when everything is linear and Gaussian.

Assumptions

- ▶ Initial state: Gaussian prior over X_0
- ▶ Dynamics model: Linear with Gaussian process noise
- ▶ Observation model: Linear with Gaussian measurement noise

Why it's useful:

- ▶ The belief at each step — the posterior $p(X_t | Y_{1:t})$ — stays Gaussian.
- ▶ All filtering steps (prediction and update) have closed-form analytical solutions.

Key Insight

Linear-Gaussian systems allow exact, recursive inference — no sampling or approximations needed.

Kalman Filter: Formal Definition

The Kalman Filter is defined by three components:

1. Initial State (Prior Belief):

$$X_0 \sim \mathcal{N}(\mu_0, \Sigma_0) \quad (3.1)$$

2. State Transition Model (Dynamics):

$$X_{t+1} = FX_t + \varepsilon_t, \quad \varepsilon_t \sim \mathcal{N}(0, \Sigma_x) \quad (3.2)$$

- F : State transition matrix ($d \times d$)
- ε_t : Process noise (may include drift; zero-mean assumed here)

3. Observation Model (Sensor):

$$Y_t = HX_t + \eta_t, \quad \eta_t \sim \mathcal{N}(0, \Sigma_y) \quad (3.3)$$

- H : Observation matrix ($m \times d$)
- η_t : Measurement noise

Note: F, H, Σ_x, Σ_y are assumed to be known (can be time-varying).

Graphical Model and Conditional Independencies

Kalman Filter as a Probabilistic Graphical Model:

- ▶ Hidden state sequence: $X_1 \rightarrow X_2 \rightarrow X_3 \rightarrow \dots$
- ▶ Each X_t emits an observation Y_t

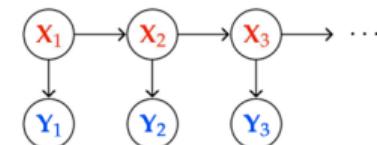


Figure 3.2: Directed graphical model of a Kalman filter with hidden states X_t and observables Y_t .

Conditional Independencies

- ▶ Markov Property (Eq. 3.4):

$$X_{t+1} \perp\!\!\!\perp X_{1:t-1}, Y_{1:t-1} \mid X_t$$

Future state depends only on the current state.

Graphical Model and Conditional Independencies

Conditional Independencies

- ▶ **Observation Dependence (Eq. 3.5):**

$$Y_t \perp\!\!\!\perp X_{1:t-1} \mid X_t$$

Observation depends only on the current state.

- ▶ **Observation Independence (Eq. 3.6):**

$$Y_t \perp\!\!\!\perp Y_{1:t-1} \mid X_{t-1}$$

Given the past state, current observation is independent of past observations.

The Goal: Online Recursive Estimation

What we want:

- ▶ At each time t , compute the belief:

$$p(X_t \mid Y_{1:t})$$

- ▶ Without re-processing all past data from scratch!

Why Online?

- ▶ Observations arrive sequentially in real-time.
- ▶ We need to update our estimate efficiently and recursively.
- ▶ No need to store all $Y_{1:t}$ — just the current posterior!

Recursive Filtering Loop

1. Predict: Use dynamics model to forecast next state.
2. Update: Incorporate new observation.
3. Repeat...

This is the essence of recursive Bayesian estimation.

Algorithm: Bayesian Filtering

Algorithm 3.2: Bayesian filtering

- 1 start with a prior over initial states $p(x_0)$
 - 2 **for** $t = 1$ **to** ∞ **do**
 - 3 **assume we have** $p(x_t | y_{1:t-1})$
 - 4 **conditioning:** compute $p(x_t | y_{1:t})$ using the new observation y_t
 - 5 **prediction:** compute $p(x_{t+1} | y_{1:t})$
-

Bayesian Filtering: Prediction and Conditioning

1. Conditioning (Update Step)

Incorporate new observation Y_t to refine belief about X_t

$$p(X_t | Y_{1:t}) \propto p(Y_t | X_t) \cdot p(X_t | Y_{1:t-1}) \quad (3.8)$$

- ▶ Combines the prior (predicted) belief with the new observation.
- ▶ Uses Bayes' Rule.

Bayesian Filtering: Prediction and Conditioning

2. Prediction (Time Update Step)

Project the belief forward to estimate X_{t+1}

$$p(X_{t+1} \mid Y_{1:t}) = \int p(X_{t+1} \mid X_t) \cdot p(X_t \mid Y_{1:t}) dX_t \quad (3.9)$$

- ▶ Uses the transition model to forecast the next state.
- ▶ Allows filtering to continue over time.

Deriving the Conditioning (Update) Step

Goal: Update belief about current state given a new observation:

$$p(X_t | Y_{1:t}) = \text{posterior}$$

Start from:

$$p(X_t | Y_{1:t}) = p(X_t | Y_t, Y_{1:t-1})$$

Apply Bayes' Rule:

$$p(X_t | Y_{1:t}) \propto p(Y_t | X_t, Y_{1:t-1}) \cdot p(X_t | Y_{1:t-1})$$

Use Conditional Independence (Eq. 3.6):

$$Y_t \perp Y_{1:t-1}, X_{1:t-1} \setminus X_t \mid X_t \Rightarrow p(Y_t | X_t, Y_{1:t-1}) = p(Y_t | X_t)$$

Final form:

$$p(X_t | Y_{1:t}) \propto p(Y_t | X_t) \cdot p(X_t | Y_{1:t-1}) \tag{3.8}$$

Interpretation

Posterior \propto Likelihood \times Prior

Deriving the Prediction Step

Goal: Predict the next state before the next observation arrives:

$$p(X_{t+1} \mid Y_{1:t})$$

Step 1: Apply marginalization

$$p(X_{t+1} \mid Y_{1:t}) = \int p(X_{t+1}, X_t \mid Y_{1:t}) dX_t$$

Step 2: Use the product rule

$$= \int p(X_{t+1} \mid X_t, Y_{1:t}) \cdot p(X_t \mid Y_{1:t}) dX_t$$

Step 3: Apply the Markov Property (Eq. 3.4):

$$X_{t+1} \perp Y_{1:t} \mid X_t \Rightarrow p(X_{t+1} \mid X_t, Y_{1:t}) = p(X_{t+1} \mid X_t)$$

Final Result:

$$p(X_{t+1} \mid Y_{1:t}) = \int p(X_{t+1} \mid X_t) \cdot p(X_t \mid Y_{1:t}) dX_t \quad (3.9)$$

Kalman Filter: The Prediction Step

Given: Posterior at time t :

$$X_t \mid Y_{1:t} \sim \mathcal{N}(\mu_t, \Sigma_t)$$

We want: Predictive belief for time $t + 1$:

$$X_{t+1} \mid Y_{1:t} \sim \mathcal{N}(\hat{\mu}_{t+1}, \hat{\Sigma}_{t+1})$$

Apply the transition model:

$$X_{t+1} = FX_t + \varepsilon_t, \quad \varepsilon_t \sim \mathcal{N}(0, \Sigma_x)$$

Predicted Mean (Eq. 3.20)

$$\hat{\mu}_{t+1} = F\mu_t$$

Predicted Covariance (Eq. 3.21)

$$\hat{\Sigma}_{t+1} = F\Sigma_t F^\top + \Sigma_x$$

Kalman Filter: The Update Step

We have:

- ▶ Prior prediction: $X_{t+1} | Y_{1:t} \sim \mathcal{N}(\hat{\mu}_{t+1}, \hat{\Sigma}_{t+1})$
- ▶ New observation: $Y_{t+1} = HX_{t+1} + \eta_{t+1}, \eta_{t+1} \sim \mathcal{N}(0, \Sigma_y)$

Kalman Gain (Eq. 3.18d)

$$K_{t+1} = \hat{\Sigma}_{t+1} H^\top \left(H \hat{\Sigma}_{t+1} H^\top + \Sigma_y \right)^{-1}$$

Updated Mean (Eq. 3.18b)

$$\mu_{t+1} = \hat{\mu}_{t+1} + K_{t+1} (Y_{t+1} - H\hat{\mu}_{t+1})$$

Updated Covariance (Eq. 3.18c)

$$\Sigma_{t+1} = (I - K_{t+1} H) \hat{\Sigma}_{t+1}$$

Example 3.4 – Random Walk in 1D

System Setup:

- **State:** $X_t \in \mathbb{R}$ (position on a line)

- **Motion Model:**

$$X_{t+1} = X_t + \varepsilon_t, \quad \varepsilon_t \sim \mathcal{N}(0, \sigma_x^2)$$

- **Sensor Model:**

$$Y_t = X_t + \eta_t, \quad \eta_t \sim \mathcal{N}(0, \sigma_y^2)$$

- **Initial Belief:** $X_0 \sim \mathcal{N}(\mu_0, \sigma_0^2)$

Kalman Filter simplifies to 1D:

- Prediction and update equations become scalar operations.
- Still recursive!
- Posterior: $X_t | Y_{1:t} \sim \mathcal{N}(\mu_t, \sigma_t^2)$

Key Idea

Even this simple setup captures the trade-off between motion uncertainty and noisy observations.

Random Walk – 1D Kalman Update (Eq. 3.13)

Recall:

- ▶ Prior at time t : $X_t \mid Y_{1:t} \sim \mathcal{N}(\mu_t, \sigma_t^2)$

- ▶ Prediction for $t + 1$:

$$X_{t+1} \mid Y_{1:t} \sim \mathcal{N}(\mu_{t+1}, \sigma_{t+1}^2)$$

- ▶ Observation: $Y_{t+1} = X_{t+1} + \eta_{t+1}$, $\eta_{t+1} \sim \mathcal{N}(0, \sigma_y^2)$

Update step (Eq. 3.13):

$$\mu_{t+1} = \frac{\sigma_y^2 \mu_t + (\sigma_t^2 + \sigma_x^2) Y_{t+1}}{\sigma_t^2 + \sigma_x^2 + \sigma_y^2}$$

$$\sigma_{t+1}^2 = \frac{(\sigma_t^2 + \sigma_x^2) \sigma_y^2}{\sigma_t^2 + \sigma_x^2 + \sigma_y^2}$$

Kalman Gain λ and Equivalent Update Forms

Kalman Gain in 1D (Eq. 3.14):

$$\lambda = \frac{\sigma_t^2 + \sigma_x^2}{\sigma_t^2 + \sigma_x^2 + \sigma_y^2}$$

Updated Mean (Eq. 3.15, 3.16):

$$\mu_{t+1} = (1 - \lambda)\mu_t + \lambda Y_{t+1} \quad (3.15)$$

$$= \mu_t + \lambda(Y_{t+1} - \mu_t) \quad (3.16)$$

Updated Variance (Eq. 3.17):

$$\sigma_{t+1}^2 = (1 - \lambda)(\sigma_t^2 + \sigma_x^2) = \lambda\sigma_y^2$$

Bayesian Linear Regression as a Kalman Filter

Problem Setup:

- We want to estimate a fixed but unknown weight vector $w^* \in \mathbb{R}^d$
- Observations arrive sequentially: (x_t, y_t) , where

$$y_t = x_t^\top w^* + \eta_t, \quad \eta_t \sim \mathcal{N}(0, \sigma_n^2)$$

Kalman Filter Interpretation:

- Hidden state: $X_t = w^*$ (constant!)
- Transition model: $X_{t+1} = X_t + 0 \rightarrow$ no process noise $\rightarrow \Sigma_x = 0$
- Sensor model:

$$Y_t = x_t^\top X_t + \eta_t$$

- Observation matrix: $H_t = x_t^\top$

Connection

Online Bayesian Linear Regression is Kalman Filtering with constant hidden state.

Bayesian Smoothing: Looking Back With More Info

Filtering: Estimates the current state using data **up to** time t

$$p(X_k \mid Y_{1:k}) \quad (\text{for } k = t)$$

Smoothing: Estimates a **past** state X_k using all data **up to a later time** t

$$p(X_k \mid Y_{1:t}) \quad \text{where } k < t$$

Why Smooth?

- ▶ Gives more accurate estimate of X_k using future evidence.
- ▶ Useful in offline settings (e.g., analyzing recorded sensor data).

Bayesian Smoothing: Looking Back With More Info

Key Equation (Eq. 3.10):

$$p(x_k \mid y_{1:t}) \propto p(x_k \mid y_{1:k}) \cdot p(y_{k+1:t} \mid x_k)$$

- ▶ First term: filtering result
- ▶ Second term: likelihood of future observations
- ▶ Smoothing = Filtering + Backward Pass
- ▶ We combine past estimate with future information for better accuracy

Kalman Smoothing: Fwd-Bkwd Estimation

Filtering: Real-time estimate

$$p(X_k \mid Y_{1:k}) = \text{filtering posterior}$$

Smoothing: Offline refinement

$$p(X_k \mid Y_{1:t}) = \text{smoothing posterior}, \quad k < t$$

Linear-Gaussian Case:

- ▶ Both filtering and smoothing posteriors are Gaussian.
- ▶ Can compute smoothing efficiently using a backward recursion.

Kalman Smoothing: Fwd-Bkwd Estimation

Key Equation (Eq. 3.11):

$$p(x_k \mid y_{1:t}) = \mathcal{N}(\tilde{\mu}_k, \tilde{\Sigma}_k)$$

- ▶ Based on forward-filtered posterior $p(x_k \mid y_{1:k})$
- ▶ And the backward Kalman gain, derived from transition model

Smoothing Algorithm = Filter Forward + Smooth Backward

Also called the Forward-Backward or Two-Filter smoother.



Thank you!



OPTIMIZATION AND ESTIMATION LAB

ONE.UNM.EDU